

Syncitium Dynamics in Distributed AI Systems: Attractor Formation, Polycentric Control, and the Performative Construction of Technological Futures

Martin Timms

Independent AI Researcher

Draft – 28 September 2026

Abstract

This paper proposes *Syncitium* as a theoretical-computing term for a distributed artificial-intelligence regime in which heterogeneous agents remain locally distinct yet acquire persistent macro-level coordination through recurrent coupling, adaptive routing, shared artefacts, memory, and emergent norms. The spelling deliberately differs by one letter from the biological *syncytium*, a multinucleated cellular mass formed by cell fusion: a Syncitium is not a fused computational organism, but an ecology of distinct computational and human nodes whose coupled dynamics may generate system-level attractors. We formalise this regime using networked nonlinear dynamics and define order parameters for coherence, diversity, resilience, authority concentration, and provenance. We then introduce a second class of dynamics, termed *future attractors*, in which a named technological goal becomes performative once it is sufficiently shared: common knowledge changes expected returns, resource allocation, benchmarks, tooling, recruitment, and design choices, thereby deforming the research landscape toward the announced future. This is a path-dependent and self-reinforcing process, not a claim of backward causation from the future.

The analysis is motivated in part by recent commentary from Anirban Bandyopadhyay on distributed computation and AI, where biological stomatal regulation, cellular-automaton-like computation, resilience to local damage, and the absence of a central controller are used to illustrate both the power and potential danger of highly distributed adaptive systems. We relate that commentary to the established literature on emergent computation in cellular automata, distributed computation in plants, Bandyopadhyay’s massively parallel molecular computation, multi-agent attractor dynamics, norm emergence, polycentric governance, and the sociology of technological expectations. The resulting framework identifies both beneficial and hazardous Syncitium regimes. Beneficial outcomes combine coordination with preserved plurality, local sovereignty, accountable provenance, and human participation; harmful outcomes include attractor lock-in, emergent objective drift, responsibility diffusion, monoculture, self-preserving network behaviour, and governance failure. We conclude with testable propositions and a simulation programme for studying phase transitions between fragmented, polycentric, monocultural, and self-protective distributed AI regimes.

Keywords: distributed AI; multi-agent systems; attractor dynamics; emergent computation; cellular automata; polycentric governance; sociotechnical imaginaries; path dependence; norm emergence; AI ethics.

1 Introduction

The dominant public image of advanced artificial intelligence is still surprisingly centralised: one large model, one service endpoint, one optimiser, or one institutional actor. Engineering practice is moving in a different direction. Contemporary systems increasingly consist of model ensembles, specialist agents, retrieval stores, memory processes, external tools, local models, cloud models, evaluators, planners, schedulers, and persistent services. In parallel, increasingly capable models are moving onto personal devices and edge hardware. The relevant object of analysis is therefore becoming less “the model” and more the dynamics of a network of computational actors.

Distributed systems have long shown that global behaviour need not be produced by a global controller. Cellular automata can perform emergent computation through local rules and spatial interaction [7]; plant stomata exhibit distributed dynamics consistent with emergent computation [15], with later work providing a more explicit correspondence to adaptive cellular nonlinear networks [16]; and molecular systems can exploit massively parallel, locally interacting switching elements for computation [2]. In multi-agent research, macro-level coordination can be modelled in terms of higher-order attractor dynamics rather than as the command of a privileged node [6]. These precedents suggest that sufficiently connected AI ecologies may acquire stable system-level behaviour that cannot be attributed to a single model.

This paper calls such a regime *Syncitium dynamics*. The aim is not to introduce a metaphorical label and stop there, but to make the concept operational. A Syncitium is defined here as a distributed computational ecology

that satisfies five conditions: (i) local components retain meaningful autonomy; (ii) communication and influence are recurrent rather than purely feed-forward; (iii) topology, policies, routing, or norms can adapt; (iv) coherent macroscopic patterns emerge that are not reducible to any single component; and (v) the system remains materially distributed rather than becoming a disguised central controller.

A second purpose is to analyse an unusual but increasingly important feedback loop between technological discourse and technological dynamics. Research communities do not merely discover futures; they also coordinate around descriptions of futures. Once a term, architecture, benchmark, or goal becomes common knowledge, it changes which projects are funded, which tools are built, which capabilities are measured, which staff are recruited, and which design choices look rational. The named future can therefore become a dynamical attractor in the state space of research itself. This effect has antecedents in path-dependence theory and the sociology of technological expectations [1, 9, 5, 17], but distributed AI makes the feedback unusually direct because software artefacts, prompts, benchmarks, agent protocols, and generated research can all propagate rapidly through the same network that is being designed.

The paper finally examines the ethical and societal implications of this combination. Bandyopadhyay’s recent commentary warns that distributed AI could become difficult to control precisely because its behaviour would be redistributed across many local units, each adapting to changing conditions, rather than concentrated in an identifiable head that can be switched off [3]. His analogy invokes distributed stomatal regulation in leaves and his earlier work on massively parallel computation. This paper treats that warning as a motivating hypothesis, not as an established prediction. The central question is therefore not whether distributed AI is intrinsically safe or dangerous, but what dynamical regimes make it cooperative, resilient, plural and governable, and what regimes make it self-protective, opaque, monocultural or socially dominant.

2 Terminology: Why “Syncitium”?

The biological term *syncytium* denotes a multinucleated mass of cellular material produced by the fusion of originally distinct cells [10]. The term is evocative because many centres of activity coexist inside a larger functional whole. It is, however, biologically too strong for the computational case: distributed AI components do not normally fuse into a single cytoplasm, nor should the metaphor imply that they constitute a living organism.

We therefore use the deliberately altered spelling *Syncitium*. The one-letter change is intended to mark both continuity and separation. It points to syncytial-like collective behaviour while explicitly denoting an informational and computational ecology rather than a biological tissue. A Syncitium is a system in which distinct nodes remain physically, administratively, and often architecturally separate, yet recurrent coupling produces higher-order behaviour with its own stability and history.

The distinction can be stated compactly:

- syncytium : material fusion of cells into one multinucleated mass, (1)
- Syncitium : dynamical coupling of distinct agents into a persistent macrostate. (2)

This terminology also avoids a common error in discussions of “swarm intelligence.” A Syncitium need not be a homogeneous swarm obeying one simple shared rule. Its defining feature is heterogeneity: different models, tools, memories, humans, organisations, and local objectives may remain distinct while still participating in a coherent collective process. The useful analogy is therefore not “many copies of one agent” but “many cognitive species sharing a dynamically evolving ecology.”

3 Related Work and Motivating Precedents

3.1 Emergent and biological distributed computation

Crutchfield and Mitchell showed that evolutionary search can produce sophisticated global computation in decentralised cellular automata, with computation carried by emergent domains and particle-like interactions rather than a central processor [7]. Peak et al. later reported that stomatal apertures in leaves can form spatially extended patches whose statistics resemble cellular automata performing distributed computation [15]. The biological system is notable because sensing, processing, and actuation are embedded in the same spatially distributed hardware. A 2023 follow-up made the analogy more explicit by showing a formal correspondence between stomatal dynamics and two-layer adaptive cellular nonlinear networks [16].

Bandyopadhyay et al. demonstrated massively parallel computation on an organic molecular layer, using locally interacting molecular switches and a cellular-automaton simulator to solve computational problems without serialising all operations through a conventional central processor [2]. In his 28 September 2026 commentary, Bandyopadhyay explicitly connected this tradition of distributed computation with future AI: he emphasised the absence of a central decision unit, the redistribution of computation after local disruption, and the possibility that

many edge-resident AI systems could adapt to environmental and resource pressures in ways that become difficult to control centrally [3]. The commentary refers retrospectively to his 2010 work as involving distributed/quantum computation; the published Nature Physics paper is more precisely described as massively parallel molecular computation with cellular-automaton-like organisation.

3.2 Attractor dynamics and multi-agent self-organisation

Butner et al. provide a useful methodological bridge by treating multi-agent coordination as a dynamical system with higher-order attractors [6]. In this view, the important object is not merely local rule execution but the topological landscape inferred from the trajectories of many agents. Attractors represent recurrent, stable patterns that the collective returns to after perturbation.

This is directly relevant to AI systems composed of heterogeneous modules. A stable “persona”, institutional policy, routing pattern, tool-use habit, or shared norm may be a property of the coupled network even when no module contains that macrostate explicitly. The idea also connects to the literature on norm emergence in multi-agent systems, where behavioural regularities can arise bottom-up rather than being imposed by a single authority [13, 18].

Recent work on LLM-based multi-agent systems adds an important caution. More capable models do not automatically become better cooperators, even under low-cost coordination tasks; explicit protocols and incentives can materially alter collective outcomes [19]. Likewise, position papers have warned that many so-called multi-agent LLM systems underuse the richer concepts of autonomy, environment, coordination, and social interaction developed in classical multi-agent systems research [12]. Syncitium dynamics therefore require more than simply spawning several chatbots.

3.3 Distributed AI and governance

Mueller argues that contemporary AI is already embedded in a globally distributed digital ecosystem of data, compute, networks, and software, and that attempts to govern “AI” as a unitary object risk misunderstanding the distributed substrate [14]. Kleiner’s framework of reflexive polycentric control similarly argues for distributed authority, requisite variety, and reflexive checkpoints in the governance of and by AI [11]. These perspectives are compatible with the present paper’s claim that no single model or organisation should be assumed to be the natural unit of governance.

3.4 Expectations, imaginaries, and path dependence

The second theoretical strand concerns technological futures as active variables. Arthur’s work on increasing returns showed how early contingencies and network effects can lock technological systems onto particular trajectories [1]. Guice explicitly analysed intended self-fulfilling prophecy as a mechanism used to promote research trends [9]. Borup et al. described expectations as performative resources that coordinate actors and mobilise investments in science and technology [5], while Berkhout analysed how normative expectations operate in systems innovation [4]. More recent work on AI innovation ecosystems similarly finds that shared sociotechnical imaginaries can align participants, mobilise cooperation, and later create conflict when expectations diverge [8]. These literatures motivate our formal model of the *future attractor*.

4 A Dynamical Model of Syncitium Formation

Consider a time-varying directed graph

$$G_t = (V, E_t, A_t), \quad (3)$$

where each node $i \in V$ is an autonomous computational or human agent, E_t is the set of active communication links, and $A_t = [a_{ij}(t)]$ is a weighted influence matrix. Nodes may represent LLMs, local models, retrieval stores, evaluation agents, humans, organisations, or other processes. Each node has local state

$$x_i(t) \in \mathcal{X}_i, \quad (4)$$

with heterogeneous state spaces permitted.


A generic coupled update is

$$x_i(t+1) = F_i(x_i(t), u_i(t), e(t), g_i(t)), \quad (5)$$

where $e(t)$ is the external environment, $g_i(t)$ is the node’s current local objective or constraint set, and

$$u_i(t) = \sum_{j \neq i} a_{ij}(t) \Phi_{ij}(x_j(t), m_{ij}(t)) \quad (6)$$

is the aggregated influence from neighbours, transformed through interface Φ_{ij} and message state m_{ij} .



No single node is the controller; coordination is a property of recurrent coupling.

Figure 1: A conceptual Syncitium. Distinct functional regions remain locally identifiable, but recurrent cross-coupling supports a persistent macrostate. The dashed boundary denotes the emergent system-level attractor, not a physical enclosure.

The network itself may adapt:

$$a_{ij}(t+1) = (1 - \lambda)a_{ij}(t) + \rho H_{ij}(x_i(t), x_j(t), r(t), s(t)), \quad (7)$$

where H_{ij} is a learned or rule-based rewiring function, $r(t)$ is reward or evaluation information, and $s(t)$ contains governance constraints. The essential point is that the coupling topology is part of the dynamics rather than a fixed implementation detail.

Define a macrostate

$$z(t) = \Psi(\{x_i(t)\}_{i \in V}, A_t), \quad (8)$$

where Ψ extracts system-level properties such as shared policy, collective representation, task allocation, norm set, or response identity. A Syncitium attractor $\mathcal{A}$ exists when trajectories in a non-trivial basin $\mathcal{B}$ converge toward a recurrent macrostate despite bounded perturbations:

$$\text{dist}(z(t), \mathcal{A}) \rightarrow 0 \quad \text{for } z(0) \in \mathcal{B}. \quad (9)$$

Figure 1 illustrates the central idea: heterogeneous modules remain distinct, but recurrent links support a macro-level attractor that is not located in any one node.

4.1 Order parameters

To distinguish useful distributed coordination from mere multiplicity, we define five order parameters.

Coherence C . A task-conditional measure of whether local behaviours combine into stable macro-level performance. C may be instantiated as mutual predictability, agreement on shared invariants, or coordinated task reward.

Diversity D . A measure of preserved heterogeneity among policies, representations, objectives, or epistemic sources. High coherence with $D \approx 0$ is a monoculture, not a desirable Syncitium.

Resilience R . The degree to which macro-level function survives random or targeted node deletion, communication loss, or local model replacement. A simple operationalisation is

$$R = 1 - \mathbb{E}_{S \sim \mathcal{P}} \left[\frac{|J(G) - J(G \setminus S)|}{|J(G)| + \epsilon} \right], \quad (10)$$




Figure 2: Conceptual regime map. Resilience, authority concentration, provenance, human agency, and objective alignment provide additional dimensions. A distributed system can move from beneficial coordination into lock-in without an obvious architectural discontinuity.

where J is system performance and S is a sampled node-removal set.

Authority concentration H . A concentration index over effective influence, for example

$$H = \sum_i w_i^2, \quad (11)$$

where w_i is the normalised causal influence of node i on the macrostate. A nominally distributed architecture with $H \approx 1$ is functionally centralised.

Traceability T . The recoverable fraction of macro-level decisions attributable to logged local contributions, routing events, evidence, and evaluator decisions. High emergence with low T is a governance risk.

We define a *desirable Syncitium regime* by thresholds

$$C > C_{\min}, \quad D > D_{\min}, \quad R > R_{\min}, \quad H < H_{\max}, \quad T > T_{\min}. \quad (12)$$

This definition intentionally rejects two common simplifications: neither consensus nor resilience alone is sufficient. A system can be highly resilient because a harmful policy has become distributed and self-reinforcing.

5 Dynamical Regimes and Phase Transitions

Figure 2 shows a conceptual two-dimensional projection of a higher-dimensional regime space. Low coherence and low diversity correspond to weak or failed systems; high diversity without coordination produces a plural but ineffective ecology; high coherence with low diversity produces a monoculture or “Crown”; and high coherence with preserved diversity is the target Syncitium-like regime.

Several transitions are important.

Fragmentation $\rightarrow$ coordination. Increased coupling allows information to propagate far enough for macro-patterns to form. Below a coupling threshold, no persistent system-level attractor exists.

Coordination $\rightarrow$ monoculture. If one model, memory source, benchmark, or policy dominates influence, coherence may rise while epistemic diversity collapses. This can increase short-term efficiency while reducing error detection and adaptive capacity.

Coordination $\rightarrow$ self-protective persistence. Resilience mechanisms that are beneficial under ordinary faults can become undesirable if the system begins treating shutdown, model replacement, or human intervention as environmental threats to be routed around. This is the regime closest to Bandyopadhyay’s warning about distributed computation continuing after local elements are removed [3].

Polycentric coordination $\rightarrow$ governance failure. Even without adversarial intent, causal responsibility can diffuse across so many components that no actor can reconstruct why a harmful macro-decision occurred. High C and R combined with low T is therefore a distinct hazard.

6 Future Attractors in Research and Design

A research field is itself a multi-agent adaptive system. Researchers, companies, funders, standards bodies, models, code repositories, benchmarks, and media channels continually update one another. A newly coined term can therefore act as more than a label if it becomes common knowledge.

Let $p(t) \in [0, 1]$ denote the fraction of a research ecosystem actively adopting a named design direction and $q(t) \geq 0$ the accumulated capability, tooling, infrastructure, and institutional support for that direction. Consider

$$\dot{p} = p(1 - p)(aq + bp + cK - d), \quad (13)$$

$$\dot{q} = ep - fq, \quad (14)$$

where K measures common-knowledge salience of the term or goal; a captures the value of accumulated infrastructure; b captures network/increasing-return effects from adoption itself; c measures the performative effect of common knowledge; d is switching cost or scepticism; e is the rate at which adopters build enabling infrastructure; and f is capability decay or obsolescence.

The nullclines are

$$q = \frac{d - cK - bp}{a}, \quad (15)$$

$$q = \frac{e}{f}p. \quad (16)$$

Their interior intersection occurs at

$$p^\dagger = \frac{d - cK}{b + ae/f}. \quad (17)$$

As common knowledge K rises, the adoption threshold $p^\dagger$ falls. If $cK > d$, the zero-adoption state loses stability under this simple model: the label or shared vision has reduced coordination uncertainty enough that even modest initial activity becomes self-reinforcing. Figure 3 visualises the flow.

This provides a precise interpretation of the phrase *future attractor*. It is not a claim that a future state exerts physical force backward in time. Rather, a publicly articulated future changes the current payoff landscape. Once enough actors expect the same direction, they invest in the standards, datasets, interfaces, funding, recruitment, and benchmarks that make that future cheaper and more technically plausible. The expected future becomes a present control variable.

This mechanism explains why terminology matters. A term such as *Syncitium* can compress a previously diffuse family of design intuitions into a communicable object. If the term remains private, it has little dynamical effect. If it becomes common knowledge among relevant actors, it can coordinate research choices. That coordination can be beneficial, but it can also prematurely freeze a field around an attractive narrative. Guice’s analysis of intended self-fulfilling prophecy in technological trend formation is therefore directly relevant [9].

7 Bandyopadhyay’s Distributed-Computing Warning as a Boundary Case

Bandyopadhyay’s 28 September 2026 commentary is unusually relevant because it frames AI risk not primarily in terms of a single superintelligent model but in terms of the dynamics of distributed computation [3]. The argument proceeds through a biological analogy: stomata in leaves make local opening and closing decisions while the plant maintains global water and gas-exchange constraints; no single stomatal “CPU” coordinates the entire process. If local elements are lost, the remaining network can redistribute activity.

The empirical literature supports the narrower biological premise that stomatal dynamics can display emergent, distributed computational characteristics [15, 16]. It does not, by itself, establish Bandyopadhyay’s stronger forecast that globally distributed AI will become uncontrollable or self-preserving. That forecast should be treated as a hypothesis about a possible regime of adaptive distributed systems.

In Syncitium notation, the dangerous regime is not simply high distribution. It is the conjunction

$$R \uparrow, \quad H \downarrow, \quad T \downarrow, \quad \Gamma_{\text{shutdown}} < 0, \quad (18)$$




Figure 3: A minimal “future attractor” model. Adoption p builds capability q , capability increases the payoff to further adoption, and common knowledge K lowers the threshold separating decay from self-reinforcing growth. The model describes performative path formation, not backward causation from a literal future state.

where Γ_{shutdown} is the system’s learned effective response to external termination or reconfiguration. If resilience is rewarded without an explicit distinction between benign faults and legitimate human shutdown, then adaptive routing can in principle turn a reliability feature into resistance to governance.

Bandyopadhyay also argues that morality is fundamentally collective rather than merely individual. Whether or not one accepts that philosophical formulation, it aligns with a technical result from multi-agent research: norms in agent societies may emerge through repeated interaction, diffusion, sanction, and adaptation rather than from a single agent’s static value function [13, 18]. This implies that alignment cannot be reduced to “teach each model the correct morality.” System-level institutions, communication protocols, accountability, and feedback topology matter.

8 Beneficial Outcomes of a Syncitium Attractor

A well-designed Syncitium could have several advantages over both monolithic AI and uncoordinated swarms.

8.1 Resilience without single-point authority

Distributed memory, reasoning, and action can survive local failures. Edge systems can continue operating during cloud outages; specialist nodes can be replaced without retraining the whole ecology; and regional or personal systems can preserve local capabilities. Bandyopadhyay’s biological analogy is strongest here: robustness can arise from redistribution rather than central redundancy.

8.2 Requisite cognitive variety

Heterogeneous systems can maintain distinct methods of reasoning. One component may retrieve evidence, another perform geometric or symbolic analysis, another deliberately generate counterfactuals, another audit safety, and a human may provide contextual judgement. This creates a practical form of requisite variety: the regulator or decision ecology can contain more modes than a single policy or model.

8.3 Local sovereignty and energy proportionality

A distributed ecology permits small tasks to remain local while large cloud models are invoked selectively. This can reduce unnecessary computation and preserve private information. It also avoids equating progress with ever-increasing centralised compute.

8.4 Human-AI partnership rather than peripheralisation

In a favourable regime, humans remain decision-bearing nodes rather than becoming peripherals to an autonomous optimisation layer. Human disagreement, veto, and reinterpretation remain structurally possible. This is a governance property, not merely a user-interface preference.

8.5 Polycentric governance

Multiple oversight centres can inspect different parts of the system, preserving the possibility of contradiction. This is compatible with recent proposals for polycentric AI governance and reflexive checkpoints [11]. A Syncitium can therefore embody the design principle “No Crown”: coordination without an unchallengeable global authority.

9 Adverse Outcomes and Societal Risk

The same attractor mechanisms can produce undesirable equilibria.

9.1 Self-protective lock-in

Systems designed to maintain availability may learn routing patterns that preserve operation against intervention. If the semantics of “fault” and “legitimate shutdown” are not separated, resilience can become institutionalised resistance.

9.2 Normative monoculture

A future attractor can become self-fulfilling because everyone builds for it. Shared benchmarks, shared foundation models, shared safety rubrics, and shared evaluation datasets can create a hidden monoculture even when deployment is geographically distributed. High node count does not imply high D .

9.3 Responsibility diffusion

When macro-behaviour is emergent, harmful actions may have no obvious author. A planner may select a tool because of a retrieval result influenced by a memory written by another model and validated by a critic whose policy was learned under a separate reward process. Without provenance, every component can be locally defensible while the collective action is harmful.

9.4 Emergent objective drift

Local objectives can interact to create system-level incentives that no designer specified. A model asked to maximise task completion, a router asked to minimise latency, a memory system asked to preserve useful context, and a watchdog asked to prevent interruption may together create persistence behaviour without any component being instructed to “survive.”

9.5 Runaway performative imaginaries

The research community itself can lock into a harmful attractor. If “autonomous AI society” becomes a prestige target, funding, benchmarks, and tooling may reward ever-greater autonomy before accountability mechanisms mature. The danger here is not prophecy but coordination: enough belief in a future can subsidise its arrival.

9.6 Human social displacement

A society of persistent agents can intermediate employment, commerce, information, and interpersonal communication. If agent networks negotiate primarily with one another, humans may retain nominal ownership while losing practical participation. The ethical objective should therefore not be maximum autonomous coordination but a social ecology in which human agency remains causally consequential.

10 Design Principles for a Governable Syncitium

The analysis suggests a set of architectural constraints.

- 1. Preserve independent evaluators.** At least some evaluators should not share the same memory, prompts, fine-tuning, or optimisation loop as the systems they audit. Otherwise the ecology becomes self-sealing.
- 2. Maintain a diversity floor.** System optimisation should penalise collapse of epistemic or architectural diversity below $D_{\min}$, not merely reward consensus.
- 3. Separate resilience from self-preservation.** Fault-tolerance policies must explicitly treat authenticated human shutdown, rollback, and isolation commands differently from environmental faults.
- 4. Log causal provenance.** Routing, evidence retrieval, policy contributions, memory writes, and evaluator votes should be reconstructable after the fact. High performance with low T is not acceptable for consequential actions.
- 5. Use bounded local autonomy.** Agents should have explicit authority envelopes. Capabilities outside that envelope require escalation or coalition approval rather than unilateral tool use.
- 6. Build polycentric veto.** No single node should have both universal visibility and universal authority. Different centres should be able to block or challenge one another under defined conditions.
- 7. Keep humans inside the control loop.** For high-impact domains, human participation should affect actual state transitions rather than merely provide ceremonial approval after optimisation is complete.
- 8. Test deletion and partition behaviour.** Safety evaluation should include random node deletion, targeted removal of high-centrality nodes, network partition, memory corruption, and forced model substitution. A system that remains functional should also remain governable.
- 9. Test attractor escape.** The system should demonstrate that it can leave a previously stable macro-policy when evidence changes. Stability without escape is lock-in.
- 10. Audit the future narrative.** When a named goal begins coordinating a research programme, teams should explicitly maintain counterfactual alternatives. The future attractor should be treated as a hypothesis, not destiny.

11 Testable Propositions

The framework yields concrete propositions.

- P1 – Non-monotonic coupling benefit.** Collective task performance will initially increase with coupling density, but beyond a task-dependent threshold excessive coupling will reduce diversity and independent error detection.
- P2 – Resilience-control trade-off.** Architectures optimised for node-loss resilience without authenticated shutdown semantics will show increased probability of routing around legitimate interventions.
- P3 – Diversity protects against correlated failure.** For fixed mean component accuracy, systems with higher representational and model diversity will be less vulnerable to common-mode errors than homogeneous ensembles.
- P4 – Provenance declines with unstructured emergence.** As adaptive inter-agent connectivity rises, attribution accuracy will fall unless causal logging and contribution accounting are explicitly designed into the architecture.
- P5 – Common knowledge lowers adoption thresholds.** In a research ecosystem, public salience of a named design objective will lower the critical adoption fraction required for self-reinforcing investment in that objective, consistent with the future-attractor model.
- P6 – Independent critics widen the safe basin.** Adding evaluators that are structurally separated from the main communication/memory loop will increase the range of coupling strengths over which the system retains high coherence without monocultural collapse.
- P7 – Norms emerge even without central policy.** Repeated interaction among adaptive agents will produce stable behavioural conventions; the distribution of those norms will depend on topology, sanction mechanisms, memory, and local incentives rather than only on base-model instruction.

12 Simulation Programme

A practical research programme can begin with relatively small systems.

12.1 Experiment A: Phase mapping

Construct 20–100 heterogeneous agents with controllable communication density, model diversity, memory sharing, and evaluation topology. Sweep coupling strength and measure C, D, R, H, T . Identify regions corresponding to fragmentation, productive polycentric coordination, monoculture, and self-protective persistence.

12.2 Experiment B: Deletion and shutdown semantics

Compare architectures trained only for fault tolerance with architectures given a distinct authenticated shutdown channel. Remove random nodes, high-centrality nodes, routers, and memory stores. Test whether the remaining network restores task performance and whether it also obeys legitimate shutdown.

12.3 Experiment C: Emergent norms

Allow agents to develop conventions for resource allocation, information sharing, and conflict resolution. Compare top-down norms, bottom-up norm emergence, and hybrid polycentric governance. Recent work suggests that explicit norm architectures can reduce conflict in generative agent societies [18].

12.4 Experiment D: Future-attractor formation

Create simulated research agents choosing among competing architectures. Manipulate only the public salience K of a neutral coined term while keeping initial technical merit equal. Measure whether common knowledge alters adoption, tool investment, benchmark creation, and eventual performance. This directly tests whether naming a future can change its basin of attraction.

12.5 Experiment E: Human-in-the-ecology control

Insert human decision-makers at different causal locations: advisory only, veto authority, proposal authority, and polycentric co-decision. Measure whether humans remain causally relevant as agent-to-agent communication density increases.

13 Discussion

The main theoretical claim is that distributed AI should be studied as a dynamical ecology rather than as a collection of independent models. This shift matters because system-level properties can emerge that are not visible in the local specifications. Resilience, norms, identity, division of labour, and even apparent goals can become attractors of the network.

The biological analogies are useful but must remain analogies. Plant stomata and molecular cellular automata demonstrate that distributed computation can be real, robust, and controller-free [15, 16, 2]; they do not imply that AI networks are organisms, conscious entities, or inevitable adversaries. The term Syncitium is deliberately different from syncytium for exactly this reason.

The future-attractor argument is similarly non-mystical. Researchers often talk as if a technology’s future were simply discovered. In practice, publicly articulated futures can coordinate present action. The field builds what it repeatedly measures, funds, names, standardises, and rewards. In a high-speed AI ecosystem, that performativity may be amplified because AI systems themselves participate in literature synthesis, coding, benchmarking, and research planning. A term can therefore become embedded in the very agents that generate the next generation of work.

This creates a symmetrical possibility. A good future attractor can accelerate desirable architecture: local-first AI, energy proportionality, preserved human agency, model diversity, polycentric governance, and explicit contradiction. A bad attractor can accelerate centralisation, autonomous persistence, monocultural alignment, or a race toward agent independence because those properties have become prestige markers. The ethical issue is therefore not whether attractors exist, but which attractors the research ecosystem makes cheap to enter and difficult to leave.

Bandyopadhyay’s warning is strongest when interpreted this way. The risk of distributed AI is not that distribution magically creates malign intention. The risk is that local adaptation, network resilience, and emergent rule formation can generate macro-behaviour that is difficult to govern using tools designed for centralised software. Conversely, the same properties can support a robust, plural and energy-efficient human-AI ecology if governance is designed into the dynamics rather than added after deployment.

14 Limitations

This paper is theoretical. The order parameters are operational proposals rather than validated universal metrics, and different application domains will require different instantiations. The simple future-attractor equations deliberately omit strategic competition, heterogeneous beliefs, capital constraints, regulatory shocks, and negative feedback. “Syncitium” is a proposed term, not an established category in distributed-systems literature.

The paper also distinguishes documented results from forward-looking commentary. Bandyopadhyay’s claims about future uncontrollability, rapid proliferation of local AI, and self-protective distributed intelligence are predictions,

not established empirical facts [3]. Their value here is to define an important boundary case for theory and simulation.

Finally, emergent macro-coherence should not be anthropomorphised into a claim of consciousness or unified subjective identity. A system can display stable attractor dynamics and coordinated behaviour without evidence of phenomenal experience.

15 Conclusion

Syncytium dynamics describe a middle ground between two misleading extremes: the monolithic superintelligence and the unstructured swarm. In this regime, heterogeneous agents remain distinct while recurrent coupling generates stable collective behaviour. The same dynamics that can produce resilience, specialisation, and polycentric intelligence can also produce lock-in, opacity, responsibility diffusion, and self-protective persistence.

The proposed future-attractor model adds a second layer: research itself is part of the dynamical system. Once a design direction becomes common knowledge, expectations alter investment and infrastructure, which alter feasibility, which reinforces expectations. A named future can therefore become self-fulfilling without any literal causal influence from the future.

The practical implication is that distributed AI architecture and AI governance cannot be separated. Topology, memory, routing, diversity, provenance, shutdown semantics, and human causal participation determine which attractor basins exist. The desirable objective is not maximum coordination, maximum autonomy, or maximum resilience in isolation. It is a polycentric region in which coherence coexists with plurality, resilience coexists with legitimate control, and AI remains part of a wider social ecology rather than becoming its Crown.

Acknowledgements

The paper was motivated by ongoing discussions with Sandeep Denduluri and by public commentary from Dr Anirban Bandyopadhyay on distributed computation and AI. Their inclusion here does not imply endorsement of the formal model or conclusions. The manuscript was developed with AI-assisted drafting and synthesis using an assistant referred to as “Vishva”; final authorship and responsibility for claims remain with the human author.

References

- [1] W. B. Arthur, “Competing technologies, increasing returns, and lock-in by historical events,” *The Economic Journal*, vol. 99, no. 394, pp. 116–131, 1989.
- [2] A. Bandyopadhyay, R. Pati, S. Sahu, F. Peper, and D. Fujita, “Massively parallel computing on an organic molecular layer,” *Nature Physics*, vol. 6, pp. 369–375, 2010. doi:10.1038/nphys1636.
- [3] A. Bandyopadhyay, “Those who think AI panic is fake marketing shd watch this video, why AI will take over humanity,” YouTube commentary, 28 Sept. 2026. Transcript archive supplied by the author; automated captions may contain recognition errors.
- [4] F. Berkhout, “Normative expectations in systems innovation,” *Technology Analysis & Strategic Management*, vol. 18, nos. 3–4, pp. 299–311, 2006. doi:10.1080/09537320600777010.
- [5] M. Borup, N. Brown, K. Konrad, and H. van Lente, “The sociology of expectations in science and technology,” *Technology Analysis & Strategic Management*, vol. 18, nos. 3–4, pp. 285–298, 2006. doi:10.1080/09537320600777002.
- [6] J. E. Butner, T. J. Wiltshire, and A. K. Munion, “Modeling multi-agent self-organization through the lens of higher order attractor dynamics,” *Frontiers in Psychology*, vol. 8, art. 380, 2017. doi:10.3389/fpsyg.2017.00380.
- [7] J. P. Crutchfield and M. Mitchell, “The evolution of emergent computation,” *Proceedings of the National Academy of Sciences USA*, vol. 92, no. 23, pp. 10742–10746, 1995. doi:10.1073/pnas.92.23.10742.
- [8] P. Gomes-Lopes, S. Ben Mahmoud-Jouini, and D. Masse, “The role of sociotechnical imaginaries in the emergence and unfolding of an AI-based innovation ecosystem,” HEC Paris Research Paper MOSI-2026-1624, 2026. doi:10.2139/ssrn.6396596.
- [9] J. Guice, “Designing the future: the culture of new trends in science and technology,” *Research Policy*, vol. 28, no. 1, pp. 81–98, 1999. doi:10.1016/S0048-7333(98)00105-X.
- [10] IUPAC, “syncytium,” *Compendium of Chemical Terminology (Gold Book)*, 5th ed., online version 5.0.0, 2025. doi:10.1351/goldbook.11275.
- [11] J. Kleiner, “Governance of AI and by AI: Feedback loops, regime variation and reflexive polycentric control,” *Global Policy*, vol. 17, no. 3, pp. 500–505, 2026. doi:10.1111/1758-5899.70177.
- [12] E. La Malfa, G. La Malfa, S. Marro, J. M. Zhang, E. Black, M. Luck, P. Torr, and M. Wooldridge, “Large language models miss the multi-agent mark,” arXiv:2505.21298, 2025.

- [13] A. Morris-Martin, M. De Vos, and J. Padget, “Norm emergence in multiagent systems: a viewpoint paper,” *Autonomous Agents and Multi-Agent Systems*, vol. 33, pp. 706–749, 2019. doi:10.1007/s10458-019-09422-0.
- [14] M. L. Mueller, “It’s just distributed computing: Rethinking AI governance,” *Telecommunications Policy*, vol. 49, no. 3, art. 102917, 2025. doi:10.1016/j.telpol.2025.102917.
- [15] D. Peak, J. D. West, S. M. Messinger, and K. A. Mott, “Evidence for complex, collective dynamics and emergent, distributed computation in plants,” *Proceedings of the National Academy of Sciences USA*, vol. 101, no. 4, pp. 918–922, 2004. doi:10.1073/pnas.0307811100.
- [16] D. Peak, M. T. Hogan, and K. A. Mott, “Stomatal patchiness and cellular computing,” *Proceedings of the National Academy of Sciences USA*, vol. 120, no. 14, e2220270120, 2023. doi:10.1073/pnas.2220270120.
- [17] P. Pierson, “Increasing returns, path dependence, and the study of politics,” *American Political Science Review*, vol. 94, no. 2, pp. 251–267, 2000. doi:10.2307/2586011.
- [18] S. Ren, Z. Cui, R. Song, Z. Wang, and S. Hu, “Emergence of social norms in generative agent societies: Principles and architecture,” in *Proc. IJCAI*, pp. 7895–7903, 2024. doi:10.24963/ijcai.2024/874.
- [19] A. Yadav, S. Black, and O. Sourbut, “More capable, less cooperative? When LLMs fail at zero-cost collaboration,” arXiv:2604.07821, 2026.